

C++ in LoXiM

- Marek Dopiera

Motto

„A humble thing but perfectly done is noble.”

Sir Royce



Agenda

- Motivation
 - The cause of poor quality of code in LoXiM
 - Do we need C++ features?
 - Theoretical introduction to smart pointers, RAII, const correctness and exceptions
 - Examples of poor code
 - Error summary and classification
 - Presentation of some new, hopefully better written code
-
-

Motivation

- We all know the slogans which are about to appear, but ...
 - Practice has proven us wrong
 - Expressing the slogans in a programming language is sometimes tricky
 - Probable discussion on the suggested solutions will certainly be worthwhile
 - Much of the content is only my opinion, so it is natural that you may not agree with it
-
-

The cause of poor quality of code in LoXiM (examples)

- Poor C++ knowledge and an attempt to write code like in C or Java
 - Assumption, that I know better what is good or bad than C++ creators
 - Unverified gossips:
 - Exceptions are extremely slow
 - References are just syntactic sugar, which make the code harder to understand
 - Constructors shouldn't throw and should be light
 - `char*` is faster and fancier than `std::string`
 - Not using constructs which make the code more readable
 - Not knowing the C++ semantics
-
-

Do we need exceptions?

- Actually, can we avoid them?

```
$ cat no-exceptions.cpp
```

```
int main()
```

```
{
```

```
    char *a = new char[100000000000];
```

```
}
```

```
g++ -fno-exceptions -o no-exceptions no-exceptions.cpp
```

```
$ ./no-exceptions
```

```
terminate called after throwing an instance of 'std::bad_alloc'
```

```
what():  St9bad_alloc
```

– No!

- What about performance?
- What about the existing code base?
- Nevertheless, I say YES

Exceptions, example of readability

```
int DescriptorInputStream::read(char *buffer, unsigned long int off, unsigned long int
length, sigset_t *sigmask, int *cancel)
{
    fd_set rfd;
    int res;
    sigset_t old_mask;
    if (status < 0 && status != IS_STATUS_TIMEDOUT)
        return status;

    FD_ZERO(&rfd);
    FD_SET(fd, &rfd);
    pthread_sigmask(SIG_SETMASK, sigmask, &old_mask);
    if (*cancel){
        status = IS_STATUS_CANCEL;
        return status;
    }
    res = pselect(fd+1, &rfd, NULL, NULL, NULL, &old_mask);
    pthread_sigmask(SIG_SETMASK, &old_mask, NULL);
    if (res == 1){
        return read(buffer, off, length);
    } else {
        if (res == 0)
            status = IS_STATUS_TIMEDOUT;
        else
            if (errno == EINTR && *cancel)
                status = IS_STATUS_CANCEL;
            else
                status = IS_STATUS_OTHERERROR;
        return status;
    }
}
```

Exceptions, example of readability

```
void FileDataStream::read(char *buf, size_t len, const sigset_t &mask,
                          const bool &cancel)
{
    while (len){
        fd_set fds;
        FD_ZERO(&fds);
        FD_SET(fd, &fds);
        int res;
        {
            Masker m(mask);
            if (cancel)
                throw OperationCancelled();
            res = pselect(fd+1, &fds, NULL, NULL, NULL,
                        &m.get_old_mask());
        }
        if (res < 0)
            throw ReadError(errno);
        if (res != 1){
            //the OS is cheating ;)
            throw ReadError(EIO);
        }
        res = ::read(fd, buf, len);
        if (res < 0)
            throw ReadError(errno);
        if (res == 0)
            throw ReadError(ENODATA);
        len -= res;
    }
}
```


Do we need C++ constructs?

- References: YES
 - Objects: YES
 - Templates: YES
 - Answer: YES
-
-

Theoretical introduction

- Exceptions
 - C++, contrary to Java, strictly specifies objects' lifetime
 - Smart pointers
 - The former allow us to program elegant RAI and exception safety
 - Const correctness
 - References
 - Templates
-
-

Examples of poor code

- The reason is not to laugh at people, but to learn by example
- The samples are collected from almost every part of LoXiM



```

bool LoximServer::LoximSession::authorize(const char *login, const char *passwd)
{
    QueryResult *qres;
    int res;
    bool correct;
    res = execute_statement("begin", &qres);
    if (res)
        return false;
    delete qres;
    res = execute_statement(("validate " + string(login) + " " +
string(passwd)).c_str(), &qres);
    if (res){
        return false;
    }
    //do the check
    if (qres->type() != QueryResult::QB00L){
        delete qres;
        return false;
    } else {
        correct = ((QueryBoolResult*)qres)->getValue();
        delete qres;
    }
    res = execute_statement("end", &qres);
    if (res)
        return false;
    delete qres;

    if (correct) {
        stats->setUserLogin(string(login));
    }
    return correct;
}

```

Problems

- Manual static cast
 - If the instance of QueryResult was a subject of multiple inheritance, it would probably cause a SEGFAULT
 - `dynamic_cast` should be used
 - Arguments should be `const string&`
 - Redundant explicit conversion to string
-
-

```
int LoximServer::prepare()
{
    socket = new
TCIPServerSocket(const_cast<char*>(hostname.c_str()), port);
    int res;
    if ((res = socket->bind()) == SOCKET_CONNECTION_STATUS_OK){
        prepared = 1;
        return 0;
    } else
        return res;
}
```

Problems

- What is the reason for the existence of this method? To comply with RAI, it should be a part of the ctor
- The `cons_cast` is ugly, but the protocol's API enforces us to do it

```
int LogIO::readString( int fileDes, string &str )
{
    int errCode;
    unsigned int n = 0;
    int len;
    char *buffer;

    if( ( errCode = readInt( fileDes, len ) ) ) return errCode;

    buffer = new char[len];

    while( n < (unsigned) len )
    {
        int size = ::read( fileDes, buffer+n, len-n );
        if( !size ) return UNEXPECTED_EOF_ERROR;
        n += size;
        if( size == -1 ) return errno;
    }

    str.clear();
    str.append(buffer, len);
    delete[] buffer;

    return 0;
}
```

Problems

- There is a memory leak if read returns an error
 - Conclusion: writing methods with many execution paths having manual memory management is error prone
 - Transferring the responsibility of memory management to the compiler by using smart pointers solves the problem
-
-

```
LockManager::LockManager() :  
err(ErrorConsole::get_instance(EC_LOCK_MANAGER))  
{  
    transaction_locks = new TransactionIdMap;  
    map_of_locks       = new DBPhysicalIdMap;  
    single_lock_id     = 0;  
    mutex = new Mutex();  
    mutex->init();  
}
```

Problems

- Why aren't the subobjects contained in LockManager?
 - Memory leak will appear, if the second allocation fails
 - Smart pointers are enough to avoid that
 - Why does the mutex have the init() method? It should be inlined in the constructor.
 - Initialization lists should be used.
-
-

```
//From Store::Buffer  
PagePointer* getPagePointer(TransactionID* tid, unsigned short  
fileID, unsigned int pageID);
```

Problems

- It is not clear whether the caller will become the owner of the returned object or not. In this case he will.
 - If `auto_ptr` was used, hardly anybody would have any doubts
 - It is not clear how long should `tid` live. In this case this is only during the execution of this method
 - If `const auto_ptr&` or a reference were used, it would have been clear
-
-

```
//From QueryParser::TreeNode  
virtual string getCard(){return this->card;}
```



Problems

- virtual inline functions seldom are a good trade-off
 - This class is always used polymorphically, so the compiler can't inline the function, however it has to increase the executable's size
 - It is probably much more desirable to return a const reference to the string instead of the value; depending on the use, it might save creating a temporary object (in current implementation it will be created, because the function will not be inlined)
 - This method is not a mutator, so it should have a const keyword
-
-

```
const int Package::serialize(char **data)
{
    PackageBufferWriter *writer=new PackageBufferWriter();
    serializeW(writer);
    int size=writer->getSize()+5;
    *data=(char*)malloc(size);
    if(!(*data))
    {
        delete writer;
        return -1;
    };
    memcpy(*data,writer->getPackageBuffer(getPackageType()),size);
    delete writer;
    return size;
}
```


Problems

- Returning a const built-in type by value doesn't make any sense, the const keyword should be removed
- The object writer should not be allocated on heap, stack is definitely big enough
- -1 is hardcoded, the user will not have any clue what it means

```
ClassGraph::~~ClassGraph() {
    vector<LogicalID*> lidsToDel;
    for(MapOfClassVertices::iterator i = classGraph.begin(); i !=
classGraph.end(); i++) {
        lidsToDel.push_back((*i).first);
        delete (*i).second;
    }
    for(MapOfInvariantVertices::iterator i = invariants.begin(); i !=
invariants.end(); i++) {
        delete (*i).second;
    }
    for(vector<LogicalID*>::iterator i = lidsToDel.begin(); i !=
lidsToDel.end(); ++i) {
        delete (*i);
    }
}
```

Problems

- This is an example of an overblown destructor
- This destructor might throw (on `push_back` operation), which is unacceptable

```
//From QueryExecutor::QueryResult  
virtual bool equal(QueryResult *r)=0;
```

Problems

- The method should be marked as const
- It should not require write access to the argument (const keyword should be added)
- It should use a reference rather than a pointer; comparing with a NULL doesn't make sense anyway

```
int LogManager::syncLog(unsigned int toLSN)
{
    if (sync_file_range(fd, 0, toLSN, SYNC_FILE_RANGE_WAIT_BEFORE |
SYNC_FILE_RANGE_WRITE | SYNC_FILE_RANGE_WAIT_AFTER) < 0)
        return ESyncLog;
#ifdef DEBUG_MODE
    debug_printf(*ec, "LogManager::syncLog up to %d done.\n",
toLSN);
#endif
    return 0;
}
```

Problems

- `sync_file_range` is Linux-specific

```
char *Client::ClientConsole::line_provider(const char *prompt)
{
    char *line = 0;
    size_t len;
    if (file){
        if ((getline(&line, &len, file) == -1) || !line){
            fclose(file);
            file = 0;
            return line_provider(prompt);
        } else {
            if (feof(file)){
                fclose(file);
                file = 0;
            }
            len = strlen(line);
            if (len > 0 && line[len - 1] == '\n')
                line[len - 1] = 0;
            return line;
        }
    }
    else {
        line = readline(prompt);
        if (line)
            add_history(line);
        return line;
    }
}
```

Problems

- getline is Linux-specific
- Basing only on the interface, it is not clear, who becomes the owner of the return value

Problems

- Presence of the file „QueryParser/Stack.h”
- Why not use `std::stack`?

```
QueryStats* getQueryStats(string key);
```



Problems

- Who becomes the owner of the result?
- There is no reason for the parameter being passed by value. A const reference would be better.
- This method is just an accessor, it should be const

```
//From Store::NamedItems  
if (buffer == NULL)  
    debug_printf(*ec, "NULL buffer!!\n");  
page_pointer = buffer->getPagePointer(tid, STORE_FILE_, (unsigned int) i);
```

Problems

- Dereference of NULL
- A perfect example of what assigning any logic to NULL values can lead to



```
#ifndef _TRANSACTION_  
#define _TRANSACTION_
```

```
[...]
```

```
using namespace Store;  
using namespace LockMgr;  
using namespace Logs;  
using namespace Errors;  
using namespace SemaphoreLib;  
using namespace TypeCheck;
```

Problems

- If „using” keyword is used in header files, namespaces stop making any sense – they don't prevent from name conflicts anymore


```
SystemStats::SystemStats(string name) {  
    this->name = name;  
}
```



Problems

- This simple method has 2 redundant allocations and deallocations!!!
- The argument should be a const reference
- _INITIALIZER lists should be used – this way the object is allocated twice

```
namespace TypeCheckConstants {  
  
#define TC_FINAL_ERROR -112  
#define TC_MDN_NAME "__MDN__"  
#define TC_MDNT_NAME "__MDNT__"  
#define TC_SUB_OBJ_NAME "subobject"  
#define TC_MDN_ATOMIC "atomic"  
[...]
```

Problems

- Defines have nothing to do with namespaces, they are evaluated during the preprocessing phase, so the namespace specification is useless
- These defines pollute the global namespace, it is wrong. Enums inside a namespace should be used. They don't impose any overhead.

Error summary and classification

- Definite errors (least interesting)
- Error proneness
- Readability
- Performance



Definite errors

- Lack of exception safety
- Memory leaks (and other resources too, we used to have semaphores leak)
- Deadlocks (actually a kind of resource)



Error proneness (how to avoid?)

- Use RAII, let the compiler take care of destroying the object at proper time – we only specify it in terms of scopes
 - Use `shared_ptr` for common ownership
 - Use references to avoid NULL values
 - Use `dynamic_cast`
 - Semantics of some operations in C++ is weird at first, watch out for them
-
-

Readability

- C++ enables to write very suggestive code, let's make use of it:
 - If you return a reference, hardly anybody will try to free it; if you return an `auto_ptr` it means that the caller becomes the owner
 - If you use `const` modifiers you get the full power of the type system in C++

Performance

- Avoid the heap!!!
 - Unless you have good reasons for it, don't pass objects by value
 - Use const values, the compiler may optimize the code better
 - Use the „virtual” keyword only where it's appropriate, it's not Java, if there is no „virtual” keyword, the compiler has the occasion to inline a function, which may have a significant effect
 - Use inline consciously, it may both speed and slow the program
-
-

Conclusion

- Using C++ features and some widely spread techniques almost certainly reduces the amount of bugs in the code
 - It also increases the readability and performance
 - If you don't pay attention to such details, those details will become poor quality of the end product
-
-

Thank you

Thank you for listening, any questions or suggestions
are welcome, the best place to post them is
loxim-devel@lists.sourceforge.net
